

# nlmrt-vignette

John C. Nash

March 2, 2016

## Background

This vignette discusses the R package `nlmrt`, that aims to provide computationally robust tools for nonlinear least squares problems. Note that R already has the `nls()` function to solve nonlinear least squares problems, and this function has a large repertoire of tools for such problems. However, it is specifically NOT indicated for problems where the residuals are small or zero. Furthermore, it frequently fails to find a solution if starting parameters are provided that are not close enough to a solution. The tools of `nlmrt` are very much intended to cope with both these issues.

The functions are also intended to provide stronger support for bounds constraints and to introduce the capability for **masks**, that is, parameters that are fixed for a given run of the function.

`nlmrt` tools generally do not return the large `nls`-style object. However, we do provide a tool `wrapnls` that will run either `nlxb` followed by a call to `nls`. The call to `nls` is adjusted to use the `port` algorithm if there are bounds constraints.

## 1 An example problem and its solution

Let us try an example initially presented by [5] and developed by [2]. This is a model for the regrowth of pasture. We set up the computation by putting the data for the problem in a data frame, and specifying the formula for the model. This can be as a formula object, but I have found that saving it as a character string seems to give fewer difficulties. Note the " " that implies "is modeled by". There must be such an element in the formula for this package (and for `nls()`). We also specify two sets of starting parameters, that is, the `ones` which is a trivial (but possibly unsuitable) start with all parameters set to 1, and `huetstart` which was suggested in [2]. Finally we load the routines in the package `nlmrt`.

```
> options(width=60)
> pastured <- data.frame(
+ time=c(9, 14, 21, 28, 42, 57, 63, 70, 79),
```

```

+ yield= c(8.93, 10.8, 18.59, 22.33, 39.35,
+         56.11, 61.73, 64.62, 67.08))
> regmod <- "yield ~ t1 - t2*exp(-exp(t3+t4*log(time)))"
> ones <- c(t1=1, t2=1, t3=1, t4=1) # all ones start
> huetstart <- c(t1=70, t2=60, t3=0, t4=1)
> require(nlmrt)

```

Let us now call the routine `nlsmnqb` (even though we are not specifying bounds). We try both starts.

```

> anmrt <- nlxb(regmod, start=ones, trace=FALSE, data=pastured)
> print(anmrt)
nlmrt class object: x
residual sumsquares = 4648.1 on 9 observations
      after 3      Jacobian and 4 function evaluations
      name      coeff      SE      tstat      pval      gradient      JSingval
t1      38.8378      NA      NA      NA      -3.039e-11      3
t2      1.00007      NA      NA      NA      -7.748e-10      1.437e-09
t3      0.998202      NA      NA      NA      1.889e-08      2.275e-16
t4      0.996049      NA      NA      NA      4.15e-08      4.933e-26

```

```

> anmrtn <- try(nlxb(regmod, start=huetstart, trace=FALSE, data=pastured))
> print(strwrap(anmrtn))
[1] "c(0.480575683702448, 0.669264006079334,"
[2] "-2.28426563497325, 0.843862687207341,"
[3] "0.734652618487168, 0.0665106492947132,"
[4] "-0.985862291968445, -0.0250879549069225,"
[5] "0.500350456693326)"
[6] "c(1, 1, 1, 1, 1, 1, 1, 1, 1, -0.981556726091092,"
[7] "-0.948171282599527, -0.869750270888722,"
[8] "-0.758399834057036, -0.484261107837448,"
[9] "-0.223408951427337, -0.149363030476145,"
[10] "-0.086933293312168, -0.0385258954067725,"
[11] "1.12712321032759, 3.11275223693953, 7.48692917929397,"
[12] "12.9373484175607, 21.6609765596453, 20.654376815193,"
[13] "17.5183401160421, 13.0985419560081, 7.73883739451332,"
[14] "2.47654281941789, 8.21473160617155, 22.7941238760067,"
[15] "43.1098907467039, 80.9615739893344, 83.5067043689986,"
[16] "72.5808432835117, 55.6490931778844, 33.8144464340477)"
[17] "44"
[18] "32"
[19] "c(69.9553722026374, 61.6818319271118,"
[20] "-9.2088020481334, 2.37778402563407)"
[21] "8.37588360361953"
[22] "c(-Inf, -Inf, -Inf, -Inf)"

```

```
[23] "c(Inf, Inf, Inf, Inf)"
[24] "integer(0)"
```

Note that the standard `nls()` of R fails to find a solution from either start.

```
> anls <- try(nls(regmod, start=ones, trace=FALSE, data=pastured))
> print(strwrap(anls))
[1] "Error in nlsModel(formula, mf, start, wts) : singular"
[2] "gradient matrix at initial parameter estimates"

> anlsx <- try(nls(regmod, start=huetstart, trace=FALSE, data=pastured))
> print(strwrap(anlsx))
[1] "Error in nls(regmod, start = huetstart, trace ="
[2] "FALSE, data = pastured) : singular gradient"
```

In both cases, the `nls()` failed with a 'singular gradient'. This implies the Jacobian is effectively singular at some point. The Levenberg-Marquardt stabilization used in `nlxb` avoids this particular issue by augmenting the Jacobian until it is non-singular. The details of this common approach may be found elsewhere [4, Algorithm 23].

There are some other tools for R that aim to solve nonlinear least squares problems. We have not yet been able to successfully use the INRA package `nls2`. This is a quite complicated package and is not installable as a regular R package using `install.packages()`. Note that there is a very different package by the same name on CRAN by Gabor Grothendieck.

## 2 The `nls` solution

We can call `nls` after getting a potential nonlinear least squares solution using `nlxb`. Package `nlmrt` has function `wrapnls` to allow this to be carried out automatically. Thus,

```
> awnls <- wrapnls(regmod, start=ones, data=pastured, control=list(rofftest=FALSE))
> print(awnls)
Nonlinear regression model
model: yield ~ t1 - t2 * exp(-exp(t3 + t4 * log(time)))
data: data
t1    t2    t3    t4
69.96 61.68 -9.21 2.38
residual sum-of-squares: 8.38

Number of iterations to convergence: 0
Achieved convergence tolerance: 7.15e-08

> cat("Note that the above is just the nls() summary result.\n")
Note that the above is just the nls() summary result.
```

### 3 Problems specified by residual functions

The model expressions in R, such as

```
yield ~ t1 - t2*exp(-exp(t3+t4*log(time)))
```

are an extremely helpful feature of the language. Moreover, they are used to compute symbolic or automatic derivatives, so we do not have to rely on numerical approximations for the Jacobian of the nonlinear least squares problem. However, there are many situations where the expression structure is not flexible enough to allow us to define our residuals, or where the construction of the residuals is simply too complicated. In such cases it is helpful to have tools that work with R functions.

Once we have an R function for the residuals, we can use the safeguarded Marquardt routine `nlfb` from package `nlmrt` or else the routine `nls.lm` from package `minpack.lm` [1]. The latter is built on the Minpack Fortran codes of [3] implemented by Kate Mullen. `nlfb` is written entirely in R, and is intended to be quite aggressive in ensuring it finds a good minimum. Thus these two approaches have somewhat different characteristics.

Let us consider a slightly different problem, called WEEDS. Here the objective is to model a set of 12 data points (density  $y$  of weeds at annual time points  $tt$ ) versus the time index. (A minor note: use of `t` rather than `tt` in R may encourage confusion with the transpose function `t()`, so I tend to avoid plain `t`.) The model suggested was a 3-parameter logistic function,

$$y_{model} = b_1 / (1 + b_2 \exp(-b_3 tt))$$

and while it is possible to use this formulation, a scaled version gives slightly better results

$$y_{model} = 100b_1 / (1 + 10b_2 \exp(-0.1b_3 tt))$$

The residuals for this latter model (in form "model" minus "data") are coded in R in the following code chunk in the function `shobbs.res`. We have also coded the Jacobian for this model as `shobbs.jac`

```
> shobbs.res <- function(x){ # scaled Hobbs weeds problem -- residual
+ # This variant uses looping
+   if(length(x) != 3) stop("hobbs.res -- parameter vector n!=3")
+   y <- c(5.308, 7.24, 9.638, 12.866, 17.069, 23.192, 31.443, 38.558, 50.156, 62.948,
+         75.995, 91.972)
+   tt <- 1:12
+   res <- 100.0*x[1]/(1+x[2]*10.*exp(-0.1*x[3]*tt)) - y
+ }
> shobbs.jac <- function(x) { # scaled Hobbs weeds problem -- Jacobian
+   jj <- matrix(0.0, 12, 3)
+   tt <- 1:12
+   yy <- exp(-0.1*x[3]*tt) # We don't need data for the Jacobian
+   zz <- 100.0/(1+10.*x[2]*yy)
+   jj[tt,1] <- zz
+   jj[tt,2] <- -0.1*x[1]*zz*zz*yy
+   jj[tt,3] <- 0.01*x[1]*zz*zz*yy*x[2]*tt
```

```
+   return(jj)
+ }
```

With package `nlmrt`, function `nlfb` can be used to estimate the parameters of the WEEDS problem as follows, where we use the naive starting point where all parameters are 1.

```
> st <- c(b1=1, b2=1, b3=1)
> ans1 <- nlfb(st, shobbs.res, shobbs.jac, trace=FALSE)
> print(ans1)
nlmrt class object: x
residual sumsquares = 2.5873 on 12 observations
after 10 Jacobian and 14 function evaluations
  name      coeff      SE      tstat      pval      gradient      JSingval
b1      1.96186      0.1131      17.35 3.166e-08 -7.327e-06      130.1
b2      4.90916      0.1688      29.08 3.282e-10 1.433e-07      6.165
b3      3.1357      0.06863     45.69 5.768e-12 1.717e-06      2.735
```

This works very well, with almost identical iterates as given by `nlxb`. (Since the algorithms are the same, this should be the case.) Note that we turn off the `trace` output. There is also the possibility of interrupting the iterations to `watch` the progress. Changing the value of `watch` in the call to `nlfb` below allows this. In this code chunk, we use an internal numerical approximation to the Jacobian.

```
> cat("No jacobian function -- use internal approximation\n")
No jacobian function -- use internal approximation
> ans1n <- nlfb(st, shobbs.res, trace=FALSE, control=list(watch=FALSE)) # NO jacfn
> print(ans1n)
nlmrt class object: x
residual sumsquares = 2.5873 on 12 observations
after 10 Jacobian and 14 function evaluations
  name      coeff      SE      tstat      pval      gradient      JSingval
b1      1.96186      0.1131      17.35 3.166e-08 -7.326e-06      130.1
b2      4.90916      0.1688      29.08 3.282e-10 1.428e-07      6.165
b3      3.1357      0.06863     45.69 5.768e-12 1.719e-06      2.735
```

Note that we could also form the sum of squares function and the gradient and use a function minimization code. The next code block shows how this is done, creating the sum of squares function and its gradient, then using the `optimx` package to call a number of minimizers simultaneously.

```
> shobbs.f <- function(x){
+   res <- shobbs.res(x)
+   as.numeric(crossprod(res))
+ }
> shobbs.g <- function(x){
+   res <- shobbs.res(x) # This is NOT efficient -- we generally have res already calculated
+   JJ <- shobbs.jac(x)
+   2.0*as.vector(crossprod(JJ,res))
+ }
```

```

> require(optimx)
> aopx <- optimx(st, shobbs.f, shobbs.g, control=list(all.methods=TRUE))
> summary(aopx)

```

	b1	b2	b3	value	fevals	gevals	niter
BFGS	1.9619	4.9092	3.1357	2.5873	41	41	NA
CG	1.9619	4.9092	3.1357	2.5873	41	41	NA
Nelder-Mead	1.9619	4.9092	3.1357	2.5873	41	41	NA
L-BFGS-B	1.9619	4.9092	3.1357	2.5873	41	41	NA
nlm	1.9619	4.9092	3.1357	2.5873	NA	NA	50
nlminb	1.9619	4.9092	3.1357	2.5873	31	29	28
spg	1.9618	4.9091	3.1357	2.5873	188	NA	150
ucminf	1.9619	4.9092	3.1357	2.5873	45	45	NA
Rcgmin	1.9619	4.9092	3.1357	2.5873	1007	175	NA
Rvmin	1.9619	4.9092	3.1357	2.5873	148	57	NA
newuoa	1.9619	4.9092	3.1357	2.5873	357	NA	NA
bobyqa	1.9619	4.9092	3.1357	2.5873	626	NA	NA
nmkb	1.9619	4.9092	3.1357	2.5873	195	NA	NA
hjb	1.9618	4.9091	3.1357	2.5873	1342	NA	19
lbfgsb3	1.9619	4.9092	3.1357	2.5873	41	41	NA

```

convcode kkt1 kkt2 xtimes
BFGS      0 TRUE TRUE 0.003
CG         0 TRUE TRUE 0.002
Nelder-Mead 0 TRUE TRUE 0.002
L-BFGS-B  0 TRUE TRUE 0.002
nlm       0 TRUE TRUE 0.002
nlminb    0 TRUE TRUE 0.001
spg       0 TRUE TRUE 0.049
ucminf    0 TRUE TRUE 0.002
Rcgmin    1 TRUE TRUE 0.030
Rvmin     0 TRUE TRUE 0.011
newuoa    0 TRUE TRUE 0.007
bobyqa    0 TRUE TRUE 0.009
nmkb      0 FALSE TRUE 0.010
hjb       0 TRUE TRUE 0.021
lbfgsb3   0 TRUE TRUE 0.027

```

```

> cat("\nNow with numerical gradient approximation or derivative free methods\n")
Now with numerical gradient approximation or derivative free methods
> aopxn <- optimx(st, shobbs.f, control=list(all.methods=TRUE))
> summary(aopxn) # no file output

```

	b1	b2	b3	value	fevals	gevals
BFGS	1.9619	4.9092	3.1357	2.5873e+00	45	45
CG	1.9619	4.9092	3.1357	2.5873e+00	45	45
Nelder-Mead	1.9619	4.9092	3.1357	2.5873e+00	45	45
L-BFGS-B	1.9619	4.9092	3.1357	2.5873e+00	45	45
nlm	1.9619	4.9092	3.1357	2.5873e+00	NA	NA
nlminb	1.9619	4.9092	3.1357	2.5873e+00	32	93
spg	1.9619	4.9091	3.1357	2.5873e+00	184	NA
ucminf	1.9619	4.9090	3.1356	2.5873e+00	34	34
Rcgmin	NA	NA	NA	8.9885e+307	NA	NA
Rvmin	1.9619	4.9092	3.1357	2.5873e+00	83	47
newuoa	1.9619	4.9092	3.1357	2.5873e+00	357	NA
bobyqa	1.9619	4.9092	3.1357	2.5873e+00	626	NA
nmkb	1.9619	4.9092	3.1357	2.5873e+00	195	NA
hjb	1.9617	4.9091	3.1358	2.5873e+00	766	NA
lbfgsb3	1.9619	4.9092	3.1357	2.5873e+00	41	41

```

niter convcode kkt1 kkt2 xtimes
BFGS      NA      0 TRUE TRUE 0.004
CG         NA      0 TRUE TRUE 0.004
Nelder-Mead NA      0 TRUE TRUE 0.003
L-BFGS-B  NA      0 TRUE TRUE 0.003
nlm       50      0 TRUE TRUE 0.002
nlminb    27      0 TRUE TRUE 0.001
spg       153     0 TRUE TRUE 0.050
ucminf    NA      0 FALSE TRUE 0.002
Rcgmin    NA     9999 NA NA 0.000

```

Rvmmmin	NA	0	TRUE	TRUE	0.009
newuoa	NA	0	TRUE	TRUE	0.007
bobyqa	NA	0	TRUE	TRUE	0.009
nmkb	NA	0	FALSE	TRUE	0.009
hjkbb	19	0	FALSE	TRUE	0.012
lbfgsb3	NA	0	TRUE	TRUE	0.044

We see that most of the minimizers work with either the analytic or approximated gradient. The 'CG' option of function `optim()` does not do very well in either case. As the author of the original step and description and then Turbo Pascal code, I can say I was never very happy with this method and replaced it recently with `Rcgmin` from the package of the same name, in the process adding the possibility of bounds or masks constraints.

## 4 Converting an expression to a function

Clearly if we have an expression, it would be nice to be able to automatically convert this to a function, if possible also getting the derivatives. Indeed, it is possible to convert an expression to a function, and there are several ways to do this (references??). In package `nlmrt` we provide the tools `model2grfun.R`, `model2jacfun.R`, `model2resfun.R`, and `model2ssfun.R` to convert a model expression to a function to compute the gradient, Jacobian, residuals or sum of squares functions respectively. We do not provide any tool for converting a function for the residuals back to an expression, as functions can use structures that are not easily expressed as R expressions.

Below are code chunks to illustrate the generation of the residual, sum of squares, Jacobian and gradient code for the Ratkowsky problem used earlier in the vignette. The commented-out first line shows how we would use one of these function generators to output the function to a file named "testresfn.R". However, it is not necessary to generate the file.

First, let us generate the residuals. We must supply the names of the parameters, and do this via the starting vector of parameters `ones`. The actual values are not needed by `model2resfun`, just the names. Other names are drawn from the variables used in the model expression `regmod`.

```
> # jres <- model2resfun(regmod, ones, funname="myxres", file="testresfn.R")
> jres <- model2resfun(regmod, ones)
> print(jres)
function (prm, yield = NULL, time = NULL)
{
  t1 <- prm[[1]]
  t2 <- prm[[2]]
  t3 <- prm[[3]]
  t4 <- prm[[4]]
  resids <- as.numeric(eval(t1 - t2 * exp(-exp(t3 + t4 * log(time))) -
    yield))
}
<environment: 0x2f86d80>

> valjres <- jres(ones, yield=pastured$yield, time=pastured$time)
> cat("valjres:")
```

```

valjres:
> print(valjres)
[1] -7.93 -9.80 -17.59 -21.33 -38.35 -55.11 -60.73 -63.62
[9] -66.08

```

Now let us also generate the Jacobian and test it using the numerical approximations from package `numDeriv`.

```

> jjac <- model2jacfun(regmod, ones)
> print(jjac)
function (prm, yield = NULL, time = NULL)
{
  t1 <- prm[[1]]
  t2 <- prm[[2]]
  t3 <- prm[[3]]
  t4 <- prm[[4]]
  localdf <- data.frame(yield, time)
  jstruc <- with(localdf, eval({
    .expr1 <- log(time)
    .expr4 <- exp(t3 + t4 * .expr1)
    .expr6 <- exp(-.expr4)
    .value <- t1 - t2 * .expr6 - yield
    .grad <- array(0, c(length(.value), 4), list(NULL, c("t1",
      "t2", "t3", "t4")))
    .grad[, "t1"] <- 1
    .grad[, "t2"] <- -.expr6
    .grad[, "t3"] <- t2 * (.expr6 * .expr4)
    .grad[, "t4"] <- t2 * (.expr6 * (.expr4 * .expr1))
    attr(.value, "gradient") <- .grad
    .value
  })))
  jacmat <- attr(jstruc, "gradient")
  return(jacmat)
}
<environment: 0x2e9a3e8>

> # Note that we now need some data!
> valjjac <- jjac(ones, yield=pastured$yield, time=pastured$time)
> cat("valjac:")
valjac:
> print(valjjac)
      t1      t2      t3      t4
[1,] 1 -2.3724e-11 5.8040e-10 1.2753e-09
[2,] 1 -2.9683e-17 1.1296e-15 2.9812e-15
[3,] 1 -1.6172e-25 9.2317e-24 2.8106e-23
[4,] 1 -8.8110e-34 6.7062e-32 2.2347e-31
[5,] 1 -2.6154e-50 2.9859e-48 1.1160e-47
[6,] 1 -5.1229e-68 7.9375e-66 3.2092e-65
[7,] 1 -4.2297e-75 7.2434e-73 3.0010e-72
[8,] 1 -2.3044e-83 4.3849e-81 1.8629e-80
[9,] 1 -5.4670e-94 1.1740e-91 5.1298e-91

> # Now compute the numerical approximation
> require(numDeriv)
> Jn <- jacobian(jres, ones, , yield=pastured$yield, time=pastured$time)
> cat("maxabsdiff=",max(abs(Jn-valjjac)), "\n")
maxabsdiff= 3.7744e-10

```

As with the WEEDS problem, we can compute the sum of squares function and the gradient.



```

> ssfn <- model2ssfun(regmod, ones) # problem getting the data attached!
> print(ssfn)
function (prm, yield = NULL, time = NULL)
{
  t1 <- prm[[1]]
  t2 <- prm[[2]]
  t3 <- prm[[3]]
  t4 <- prm[[4]]
  resids <- as.numeric(eval(t1 - t2 * exp(-exp(t3 + t4 * log(time))) -
    yield))
  ss <- as.numeric(crossprod(resids))
}
<environment: 0x24ba310>

> valss <- ssfn(ones, yield=pastured$yield, time=pastured$time)
> cat("valss: ",valss,"\n")
valss: 17533

> grfn <- model2grfun(regmod, ones) # problem getting the data attached!
> print(grfn)
function (prm, yield = NULL, time = NULL)
{
  t1 <- prm[[1]]
  t2 <- prm[[2]]
  t3 <- prm[[3]]
  t4 <- prm[[4]]
  localdf <- data.frame(yield, time)
  jstruc <- with(localdf, eval({
    .expr1 <- log(time)
    .expr4 <- exp(t3 + t4 * .expr1)
    .expr6 <- exp(-.expr4)
    .value <- t1 - t2 * .expr6 - yield
    .grad <- array(0, c(length(.value), 4), list(NULL, c("t1",
      "t2", "t3", "t4")))
    .grad[, "t1"] <- 1
    .grad[, "t2"] <- -.expr6
    .grad[, "t3"] <- t2 * (.expr6 * .expr4)
    .grad[, "t4"] <- t2 * (.expr6 * (.expr4 * .expr1))
    attr(.value, "gradient") <- .grad
    .value
  })))
  jacmat <- attr(jstruc, "gradient")
  resids <- as.numeric(eval(t1 - t2 * exp(-exp(t3 + t4 * log(time))) -
    yield))
  grj <- as.vector(2 * crossprod(jacmat, resids))
}
<environment: 0x2047e58>

> valgr <- grfn(ones, yield=pastured$yield, time=pastured$time)
> cat("valgr:")
valgr:
> print(valgr)
[1] -6.8108e+02 3.7626e-10 -9.2051e-09 -2.0226e-08

> gn <- grad(ssfn, ones, yield=pastured$yield, time=pastured$time)
> cat("maxabsdiff=",max(abs(gn-valgr)), "\n")
maxabsdiff= 1.4376e-07

```

Moreover, we can use the Huet starting parameters as a double check on our conversion of the expression to various optimization-style functions.

```

> cat("\n\nHuetstart:")
Huetstart:

```

```

> print(huetstart)
t1 t2 t3 t4
70 60 0 1

> valjres <- jres(huetstart, yield=pastured$yield, time=pastured$time)
> cat("valjres:")
valjres:
> print(valjres)
[1] 61.063 59.200 51.410 47.670 30.650 13.890 8.270 5.380
[9] 2.920

> valss <- ssfn(huetstart, yield=pastured$yield, time=pastured$time)
> cat("valss:", valss, "\n")
valss: 13387

> valjjac <- jjac(huetstart, yield=pastured$yield, time=pastured$time)
> cat("valjac:")
valjac:
> print(valjjac)
      t1      t2      t3      t4
[1,] 1 -1.2341e-04 6.6641e-02 1.4643e-01
[2,] 1 -8.3153e-07 6.9848e-04 1.8433e-03
[3,] 1 -7.5826e-10 9.5540e-07 2.9087e-06
[4,] 1 -6.9144e-13 1.1616e-09 3.8708e-09
[5,] 1 -5.7495e-19 1.4489e-15 5.4154e-15
[6,] 1 -1.7588e-25 6.0151e-22 2.4319e-21
[7,] 1 -4.3596e-28 1.6479e-24 6.8276e-24
[8,] 1 -3.9754e-31 1.6697e-27 7.0937e-27
[9,] 1 -4.9061e-35 2.3255e-31 1.0161e-30

> Jn <- jacobian(jres, huetstart, , yield=pastured$yield, time=pastured$time)
> cat("maxabsdiff=",max(abs(Jn-valjjac)),"\n")
maxabsdiff= 5.3945e-10

> valgr <- grfn(huetstart, yield=pastured$yield, time=pastured$time)
> cat("valgr:")
valgr:
> print(valgr)
[1] 560.90509 -0.01517 8.22138 18.10084

> gn <- grad(ssfn, huetstart, yield=pastured$yield, time=pastured$time)
> cat("maxabsdiff=",max(abs(gn-valgr)),"\n")
maxabsdiff= 4.7636e-08

```

Now that we have these functions, let us apply them with `nlfb`.

```

> cat("All ones to start\n")
All ones to start

> anlfb <- nlfb(ones, jres, jjac, trace=FALSE, yield=pastured$yield, time=pastured$time)
> print(strwrap(anlfb))
[1] "c(29.90777777777472, 28.0377777777778,"
[2] "20.2477777777778, 16.5077777777778,"
[3] "-0.512222222222185, -17.2722222222222,"
[4] "-22.8922222222222, -25.7822222222222,"
[5] "-28.2422222222222)"
[6] "c(1, 1, 1, 1, 1, 1, 1, 1, 1, -2.5904803198541e-11,"
[7] "-3.48177832682637e-17, -2.11977026263411e-25,"
[8] "-1.30186504324169e-33, -5.00031799754126e-50,"
[9] "-1.31593194786314e-67, -1.22799624106577e-74,"

```

```

[10] "-7.73807145709702e-83, -2.22164181475742e-93,"
[11] "6.31486798421711e-10, 1.31950211229055e-15,"
[12] "1.20434229540269e-23, 9.85816788938708e-32,"
[13] "5.67649212680981e-48, 2.02656873886703e-65,"
[14] "2.0899294411938e-72, 1.4630618491745e-80,"
[15] "4.73981859865953e-91, 1.38751831375555e-09,"
[16] "3.48224172088318e-15, 3.66664714105283e-23,"
[17] "3.284943150308e-31, 2.12168521608018e-47,"
[18] "8.1935213090302e-65, 8.65885924351809e-72,"
[19] "6.21581130504181e-80, 2.07103901969961e-90)"
[20] "4"
[21] "3"
[22] "c(38.8377777777778, 1.00007369903129,"
[23] "0.998201661261902, 0.996048644398237)"
[24] "4648.06335555373"
[25] "c(-Inf, -Inf, -Inf, -Inf)"
[26] "c(Inf, Inf, Inf, Inf)"
[27] "NULL"

> cat("Huet start\n")
Huet start

> anlfbh <- nlfb(huetstart, jres, jjac, trace=FALSE, yield=pastured$yield, time=pastured$time)
> print(strwrap(anlfbh))
[1] "c(0.480575683702348, 0.669264006079271,"
[2] "-2.28426563497321, 0.843862687207519,"
[3] "0.734652618487608, 0.066510649295175,"
[4] "-0.985862291968068, -0.0250879549066383,"
[5] "0.500350456693454)"
[6] "c(1, 1, 1, 1, 1, 1, 1, 1, -0.981556726091092,"
[7] "-0.948171282599528, -0.869750270888723,"
[8] "-0.758399834057038, -0.484261107837453,"
[9] "-0.223408951427342, -0.14936303047615,"
[10] "-0.0869332933121718, -0.0385258954067749,"
[11] "1.12712321032759, 3.11275223693951, 7.48692917929391,"
[12] "12.9373484175606, 21.6609765596452, 20.6543768151932,"
[13] "17.5183401160423, 13.0985419560085, 7.73883739451366,"
[14] "2.47654281941788, 8.2147316061715, 22.7941238760065,"
[15] "43.1098907467036, 80.9615739893341, 83.5067043689992,"
[16] "72.5808432835129, 55.6490931778859, 33.8144464340492)"
[17] "44"
[18] "32"
[19] "c(69.9553722026373, 61.6818319271118,"
[20] "-9.2088020481334, 2.37778402563408)"
[21] "8.37588360361957"
[22] "c(-Inf, -Inf, -Inf, -Inf)"
[23] "c(Inf, Inf, Inf, Inf)"
[24] "NULL"

```

## 5 Using bounds and masks

The manual for `nls()` tells us that bounds are restricted to the 'port' algorithm.

```

lower, upper: vectors of lower and upper bounds, replicated to be as
              long as 'start'. If unspecified, all parameters are assumed
              to be unconstrained. Bounds can only be used with the
              "port" algorithm. They are ignored, with a warning, if
              given for other algorithms.

```

Later in the manual, there is the discomfoting warning:

The 'algorithm = "port"' code appears unfinished, and does not even check that the starting value is within the bounds. Use with caution, especially where bounds are supplied.

We will base the rest of this discussion on the examples in man/nlmrt-package.Rd, and use an unscaled version of the WEEDS problem.

First, let us estimate the model with no constraints.

```
> require(nlmrt)
> # Data for Hobbs problem
> ydat <- c(5.308, 7.24, 9.638, 12.866, 17.069, 23.192, 31.443,
+          38.558, 50.156, 62.948, 75.995, 91.972)
> tdat <- 1:length(ydat)
> weeddata1 <- data.frame(y=ydat, tt=tdat)
> start1 <- c(b1=1, b2=1, b3=1) # name parameters for nlxb, nls, wrapnls.
> eunsc <- y ~ b1/(1+b2*exp(-b3*tt))
> anlxb1 <- try(nlxb(eunsc, start=start1, data=weeddata1))
> print(anlxb1)
```

nlmrt class object: x

residual sumsqares = 2.5873 on 12 observations		after 18 Jacobian and 25 function evaluations					
name	coeff	SE	tstat	pval	gradient	JSingval	
b1	196.186	11.31	17.35	3.164e-08	-2.663e-07	1011	
b2	49.0916	1.688	29.08	3.281e-10	1.59e-07	0.4605	
b3	0.31357	0.006863	45.69	5.768e-12	-5.531e-05	0.04715	

Now let us see if we can apply bounds. Note that we name the parameters in the vectors for the bounds. First we apply bounds that are NOT active at the unconstrained solution.

```
> # WITH BOUNDS
> startf1 <- c(b1=1, b2=1, b3=.1) # a feasible start when b3 <= 0.25
> anlxb1 <- try(nlxb(eunsc, start=startf1, lower=c(b1=0, b2=0, b3=0),
+               upper=c(b1=500, b2=100, b3=5), data=weeddata1))
> print(anlxb1)
```

nlmrt class object: x

residual sumsqares = 2.5873 on 12 observations		after 13 Jacobian and 18 function evaluations					
name	coeff	SE	tstat	pval	gradient	JSingval	
b1	196.186	11.31	17.35	3.164e-08	-2.662e-07	1011	
b2	49.0916	1.688	29.08	3.281e-10	1.648e-07	0.4605	
b3	0.31357	0.006863	45.69	5.768e-12	-5.872e-05	0.04715	

We note that nls() also solves this case.

```
> anlsb1 <- try(nls(eunsc, start=startf1, lower=c(b1=0, b2=0, b3=0),
+               upper=c(b1=500, b2=100, b3=5), data=weeddata1, algorithm='port'))
> print(anlsb1)
```

Nonlinear regression model

```
model: y ~ b1/(1 + b2 * exp(-b3 * tt))
data: weeddata1
      b1      b2      b3
196.186 49.092  0.314
residual sum-of-squares: 2.59
```

Algorithm "port", convergence message: relative convergence (4)

Now we will change the bounds so the start is infeasible.

```
> ## Uncon solution has bounds ACTIVE. Infeasible start
> anlxb2i <- try(nlxb(eunsc, start=start1, lower=c(b1=0, b2=0, b3=0),
+               upper=c(b1=500, b2=100, b3=.25), data=weeddata1))
> print(anlxb2i)
[1] "Error in nlxb(eunsc, start = start1, lower = c(b1 = 0, b2 = 0, b3 = 0), : \n Infeasible start\n"
attr(,"class")
[1] "try-error"
attr(,"condition")
<simpleError in nlxb(eunsc, start = start1, lower = c(b1 = 0, b2 = 0, b3 = 0), upper = c(b1 = 500, b2 = 100, b3 = 0.25), data = weeddata1) : Infeasible start\n"
> anlsb2i <- try(nls(eunsc, start=start1, lower=c(b1=0, b2=0, b3=0),
+               upper=c(b1=500, b2=100, b3=.25), data=weeddata1, algorithm='port'))
> print(anlsb2i)
[1] "Error in nls(eunsc, start = start1, lower = c(b1 = 0, b2 = 0, b3 = 0), : \n Convergence failure: initial par violates constraints\n"
attr(,"class")
[1] "try-error"
attr(,"condition")
<simpleError in nls(eunsc, start = start1, lower = c(b1 = 0, b2 = 0, b3 = 0), upper = c(b1 = 500, b2 = 100, b3 = 0.25), data = weeddata1) : Convergence failure: initial par violates constraints\n"
>
```

Both `nlxb()` and `nls()` (with 'port') do the right thing and refuse to proceed. There is a minor "glitch" in the output processing of both `knitr` and `Sweave` here. Let us start them off properly and see what they accomplish.

```
> ## Uncon solution has bounds ACTIVE. Feasible start
> anlxb2f <- try(nlxb(eunsc, start=startf1, lower=c(b1=0, b2=0, b3=0),
+               upper=c(b1=500, b2=100, b3=.25), data=weeddata1))
> print(anlxb2f)
nlmrt class object: x
residual sum-squares = 29.993 on 12 observations
after 13 Jacobian and 18 function evaluations
  name      coeff      SE      tstat      pval      gradient      JSingval
b1          500U         NA         NA         NA         0          1.529
b2      87.9425         NA         NA         NA    -1.808e-10         0
b3       0.25U         NA         NA         NA         0          0
> anlsb2f <- try(nls(eunsc, start=startf1, lower=c(b1=0, b2=0, b3=0),
+               upper=c(b1=500, b2=100, b3=.25), data=weeddata1, algorithm='port'))
> print(anlsb2f)
Nonlinear regression model
model: y ~ b1/(1 + b2 * exp(-b3 * tt))
data: weeddata1
      b1      b2      b3
500.00 87.94  0.25
residual sum-of-squares: 30

Algorithm "port", convergence message: both X-convergence and relative convergence (5)
>
```

Both methods get essentially the same answer for the bounded problem, and this solution has parameters `b1` and `b3` at their upper bounds. The Jacobian elements for these parameters are zero as returned by `nlxb()`.

Let us now turn to **masks**, which functions from `nlmrt` are designed to handle. Masks are also available with packages `Rcgmin` and `Rvmmmin`. I would like to hear if other packages offer this capability.

```

> ## TEST MASKS
> anlsmnqm <- try(nlxb(eunsc, start=start1, lower=c(b1=0, b2=0, b3=0),
+   upper=c(b1=500, b2=100, b3=5), masked=c("b2"), data=weeddata1))
> print(anlsmnqm) # b2 masked
nlmrt class object: x
residual sumsquares = 6181.2 on 12 observations
after 22 Jacobian and 35 function evaluations

```

name	coeff	SE	tstat	pval	gradient	JSingval
b1	50.4013	NA	NA	NA	-0.001511	162.1
b2	1 M	NA	NA	NA	0	0.4918
b3	0.19862	NA	NA	NA	-0.0468	0

```

> anlqm3 <- try(nlxb(eunsc, start=start1, data=weeddata1, masked=c("b3")))
> print(anlqm3) # b3 masked
nlmrt class object: x
residual sumsquares = 1031 on 12 observations
after 17 Jacobian and 18 function evaluations

```

name	coeff	SE	tstat	pval	gradient	JSingval
b1	78.5698	NA	NA	NA	8.489e-08	1.944
b2	2293.71	NA	NA	NA	-1.757e-09	0.01097
b3	1 M	NA	NA	NA	0	0

```

> # Note that the parameters are put in out of order to test code.
> anlqm123 <- try(nlxb(eunsc, start=start1, data=weeddata1, masked=c("b2","b1","b3")))
> print(anlqm123) # ALL masked - fails!!
[1] "Error in nlxb(eunsc, start = start1, data = weeddata1, masked = c(\"b2\", : \n All parameters are masked\n"
attr(,"class")
[1] "try-error"
attr(,"condition")
<simpleError in nlxb(eunsc, start = start1, data = weeddata1, masked = c("b2", "b1", "b3")): All parameters are masked>

```

Finally (for nlxb) we combine the bounds and mask.

```

> ## BOUNDS and MASK
> anlqbm2 <- try(nlxb(eunsc, start=startf1, data=weeddata1,
+   lower=c(0,0,0), upper=c(200, 60, .3), masked=c("b2")))
> print(anlqbm2)
nlmrt class object: x
residual sumsquares = 6181.2 on 12 observations
after 17 Jacobian and 28 function evaluations

```

name	coeff	SE	tstat	pval	gradient	JSingval
b1	50.4016	NA	NA	NA	0.0004618	162.2
b2	1 M	NA	NA	NA	0	0.4918
b3	0.198618	NA	NA	NA	-0.0746	0

```

> anlqbm2x <- try(nlxb(eunsc, start=startf1, data=weeddata1,
+   lower=c(0,0,0), upper=c(48, 60, .3), masked=c("b2")))
> print(anlqbm2x)
nlmrt class object: x
residual sumsquares = 6206.1 on 12 observations
after 11 Jacobian and 20 function evaluations

```

name	coeff	SE	tstat	pval	gradient	JSingval
b1	48U	NA	NA	NA	0	141.2
b2	1 M	NA	NA	NA	0	0
b3	0.215971	NA	NA	NA	-0.1502	0

Turning to the function-based nlfb,

```

> hobbs.res <- function(x){ # Hobbs weeds problem -- residual
+   if(length(x) != 3) stop("hobbs.res -- parameter vector n!=3")
+   y <- c(5.308, 7.24, 9.638, 12.866, 17.069, 23.192, 31.443, 38.558, 50.156, 62.948,
+         75.995, 91.972)
+   tt <- 1:12
+   res <- x[1]/(1+x[2]*exp(-x[3]*tt)) - y
+ }
> hobbs.jac <- function(x) { # Hobbs weeds problem -- Jacobian
+   jj <- matrix(0.0, 12, 3)
+   tt <- 1:12
+   yy <- exp(-x[3]*tt)
+   zz <- 1.0/(1+x[2]*yy)
+   jj[tt,1] <- zz
+   jj[tt,2] <- -x[1]*zz*zz*yy
+   jj[tt,3] <- x[1]*zz*zz*yy*x[2]*tt
+   return(jj)
+ }
> # Check unconstrained
> ans1 <- nlfb(start1, hobbs.res, hobbs.jac)
> ans1
nlmrt class object: x
residual sumsquares = 2.5873 on 12 observations
after 18 Jacobian and 25 function evaluations

```

name	coeff	SE	tstat	pval	gradient	JSingval
b1	196.186	11.31	17.35	3.164e-08	-2.663e-07	1011
b2	49.0916	1.688	29.08	3.281e-10	1.59e-07	0.4605
b3	0.31357	0.006863	45.69	5.768e-12	-5.531e-05	0.04715

```

> ## No jacobian - use internal approximation
> ans1n <- nlfb(start1, hobbs.res)
> ans1n
nlmrt class object: x
residual sumsquares = 2.5873 on 12 observations
after 18 Jacobian and 25 function evaluations

```

name	coeff	SE	tstat	pval	gradient	JSingval
b1	196.186	11.31	17.35	3.164e-08	-2.662e-07	1011
b2	49.0916	1.688	29.08	3.281e-10	1.589e-07	0.4605
b3	0.31357	0.006863	45.69	5.768e-12	-5.526e-05	0.04715

```

> # Bounds -- infeasible start
> ans2i <- try(nlfb(start1, hobbs.res, hobbs.jac,
+   lower=c(b1=0, b2=0, b3=0), upper=c(b1=500, b2=100, b3=.25)))
> ans2i
[1] "Error in nlfb(start1, hobbs.res, hobbs.jac, lower = c(b1 = 0, b2 = 0, : \n Infeasible start\n"
attr(,"class")
[1] "try-error"
attr(,"condition")
<simpleError in nlfb(start1, hobbs.res, hobbs.jac, lower = c(b1 = 0, b2 = 0, b3 = 0), upper = c(b1 = 500, b2 = 100, b3 = 0.25)
> # Bounds -- feasible start
> ans2f <- nlfb(startf1, hobbs.res, hobbs.jac,
+   lower=c(b1=0, b2=0, b3=0), upper=c(b1=500, b2=100, b3=.25))
> ans2f
nlmrt class object: x
residual sumsquares = 29.993 on 12 observations

```

```

      after 13      Jacobian and 18 function evaluations
name      coeff      SE      tstat      pval      gradient      JSingval
b1              500U      NA      NA      NA      0      1.529
b2      87.9425      NA      NA      NA      -1.809e-10      0
b3      0.25U      NA      NA      NA      0      0

> # Mask b2
> ansm2 <- nlfb(start1, hobbs.res, hobbs.jac, maskidx=c(2))
> ansm2
nlmrt class object: x
residual sumsquares = 6181.2 on 12 observations
      after 24      Jacobian and 38 function evaluations
name      coeff      SE      tstat      pval      gradient      JSingval
b1      50.4022      NA      NA      NA      0.001528      162.2
b2      1 M      NA      NA      NA      0      0.4918
b3      0.198611      NA      NA      NA      0.04544      0

> # Mask b3
> ansm3 <- nlfb(start1, hobbs.res, hobbs.jac, maskidx=c(3))
> ansm3
nlmrt class object: x
residual sumsquares = 1031 on 12 observations
      after 17      Jacobian and 18 function evaluations
name      coeff      SE      tstat      pval      gradient      JSingval
b1      78.5698      NA      NA      NA      8.489e-08      1.944
b2      2293.71      NA      NA      NA      -1.757e-09      0.01097
b3      1 M      NA      NA      NA      0      0

> # Mask all -- should fail
> ansma <- try(nlfb(start1, hobbs.res, hobbs.jac, maskidx=c(3,1,2)))
> ansma
[1] "Error in nlfb(start1, hobbs.res, hobbs.jac, maskidx = c(3, 1, 2)) : \n All parameters are masked\n"
attr(,"class")
[1] "try-error"
attr(,"condition")
<simpleError in nlfb(start1, hobbs.res, hobbs.jac, maskidx = c(3, 1, 2)): All parameters are masked>

> # Bounds and mask
> ansmbm2 <- nlfb(startf1, hobbs.res, hobbs.jac, maskidx=c(2),
+               lower=c(0,0,0), upper=c(200, 60, .3))
> ansmbm2
nlmrt class object: x
residual sumsquares = 6181.2 on 12 observations
      after 17      Jacobian and 28 function evaluations
name      coeff      SE      tstat      pval      gradient      JSingval
b1      50.4016      NA      NA      NA      0.0004618      162.2
b2      1 M      NA      NA      NA      0      0.4918
b3      0.198618      NA      NA      NA      -0.0746      0

> # Active bound
> ansmbm2x <- nlfb(startf1, hobbs.res, hobbs.jac, maskidx=c(2),
+               lower=c(0,0,0), upper=c(48, 60, .3))
> ansmbm2x
nlmrt class object: x
residual sumsquares = 6206.1 on 12 observations
      after 11      Jacobian and 20 function evaluations
name      coeff      SE      tstat      pval      gradient      JSingval
b1      48U      NA      NA      NA      0      141.2
b2      1 M      NA      NA      NA      0      0
b3      0.215971      NA      NA      NA      -0.1502      0

```

The results match those of `nlxb()`



Finally, let us check the results above with `Rvmmmin` and `Rcgmin`. Note that this vignette cannot be created on systems that lack these codes.

```
> require(Rcgmin)
> require(Rvmmmin)
> hobbs.f <- function(x) {
+   res<-hobbs.res(x)
+   as.numeric(crossprod(res))
+ }
> hobbs.g <- function(x) {
+   res <- hobbs.res(x) # Probably already available
+   JJ <- hobbs.jac(x)
+   2.0*as.numeric(crossprod(JJ, res))
+ }
> # Check unconstrained
> a1cg <- Rcgmin(start1, hobbs.f, hobbs.g)
> a1cg
$par
      b1      b2      b3
196.18475 49.09147 0.31357

$value
[1] 2.5873

$counts
[1] 1011 197

$convergence
[1] 1

$message
[1] "Too many function evaluations (> 1000) "
> a1vm <- Rvmmmin(start1, hobbs.f, hobbs.g)
> a1vm
$par
      b1      b2      b3
196.18626 49.09164 0.31357

$value
[1] 2.5873

$counts
function gradient
      215      55

$convergence
[1] 3

$message
[1] "Rvmmminu appears to have converged"
> ## No jacobian - use internal approximation
> a1cgn <- try(Rcgmin(start1, hobbs.f))
> a1cgn
$par
      b1      b2      b3
196.19292 49.09238 0.31357
```

```

$value
[1] 2.5873

$counts
[1] 1009 238

$convergence
[1] 1

$message
[1] "Too many function evaluations (> 1000) "

  > alvmn <- try(Rvmmmin(start1, hobbs.f))
  > alvmn

$par
      b1      b2      b3
196.18634 49.09163 0.31357

$value
[1] 2.5873

$counts
function gradient
      206      50

$convergence
[1] 3

$message
[1] "Rvmmminu appears to have converged"

  > # But
  > grfwd <- function(par, userfn, fbase=NULL, eps=1.0e-7, ...) {
+   # Forward different gradient approximation
+   if (is.null(fbase)) fbase <- userfn(par, ...) # ensure we function value at par
+   df <- rep(NA, length(par))
+   teps <- eps * (abs(par) + eps)
+   for (i in 1:length(par)) {
+     dx <- par
+     dx[i] <- dx[i] + teps[i]
+     df[i] <- (userfn(dx, ...) - fbase)/teps[i]
+   }
+   df
+ }

  > alvmn <- try(Rvmmmin(start1, hobbs.f, gr="grfwd"))
  > alvmn

$par
      b1      b2      b3
196.18634 49.09163 0.31357

$value
[1] 2.5873

$counts
function gradient
      206      50

$convergence
[1] 3

$message

```

```

[1] "Rvminu appears to have converged"

> # Bounds -- infeasible start
> # Note: These codes move start to nearest bound
> a1cg2i <- Rcgmin(start1, hobbs.f, hobbs.g,
+   lower=c(b1=0, b2=0, b3=0), upper=c(b1=500, b2=100, b3=.25))
> a1cg2i

$par
      b1      b2      b3
500.000  87.942   0.250

$value
[1] 29.993

$counts
[1] 87 45

$convergence
[1] 0

$message
[1] "Rcgmin seems to have converged"

$bdmsk
[1] -1  1 -1

> alvm2i <- Rvmin(start1, hobbs.f, hobbs.g,
+   lower=c(b1=0, b2=0, b3=0), upper=c(b1=500, b2=100, b3=.25))
> alvm2i # Fails to get to solution!

$par
      b1      b2      b3
35.532   0.000   0.250

$value
[1] 9205.4

$counts
function gradient
      6          4

$convergence
[1] 2

$message
[1] "Rvminb appears to have converged"

$bdmsk
[1]  1 -3 -1

> # Bounds -- feasible start
> a1cg2f <- Rcgmin(startf1, hobbs.f, hobbs.g,
+   lower=c(b1=0, b2=0, b3=0), upper=c(b1=500, b2=100, b3=.25))
> a1cg2f

$par
      b1      b2      b3
500.000  87.942   0.250

$value
[1] 29.993

$counts
[1] 67 34

$convergence

```

```

[1] 0

$message
[1] "Rcgmin seems to have converged"

$bdmsk
[1] -1 1 -1

> a1vm2f <- Rvmmmin(startf1, hobbs.f, hobbs.g,
+   lower=c(b1=0, b2=0, b3=0), upper=c(b1=500, b2=100, b3=.25))
> a1vm2f # Gets there, but only just!

$par
      b1      b2      b3
3.5532e+01 1.6777e-15 2.5000e-01

$value
[1] 9205.4

$counts
function gradient
      31         5

$convergence
[1] 2

$message
[1] "Rvmmminb appears to have converged"

$bdmsk
[1] 1 -3 -1

> # Mask b2
> a1cgm2 <- Rcgmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,0,1))
> a1cgm2

$par
      b1      b2      b3
50.40179 1.00000 0.19861

$value
[1] 6181.2

$counts
[1] 1006 129

$convergence
[1] 1

$message
[1] "Too many function evaluations (> 1000) "

$bdmsk
[1] 1 0 1

> a1vmm2 <- Rvmmmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,0,1))
> a1vmm2

$par
      b1      b2      b3
50.40179 1.00000 0.19861

$value
[1] 6181.2

$counts
function gradient
      290       28

```

```

$convergence
[1] 3

$message
[1] "Rvmmminb appears to have converged"

$bdmsk
[1] 1 0 1

> # Mask b3
> a1cgm3 <- Rcgmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,1,0))
> a1cgm3

$par
      b1      b2      b3
78.571 2293.937  1.000

$value
[1] 1031

$counts
[1] 172  71

$convergence
[1] 0

$message
[1] "Rcgmin seems to have converged"

$bdmsk
[1] 1 1 0

> a1vmm3 <- Rvmmmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,1,0))
> a1vmm3

$par
      b1      b2      b3
78.571 2293.947  1.000

$value
[1] 1031

$counts
function gradient
      88      30

$convergence
[1] 0

$message
[1] "Rvmmminb appears to have converged"

$bdmsk
[1] 1 1 0

> # Mask all -- should fail
> a1cgma <- Rcgmin(start1, hobbs.f, hobbs.g, bdmsk=c(0,0,0))
> a1cgma

$par
b1 b2 b3
1  1  1

$value
[1] 23521

$counts
[1] 1 1

```

```

$convergence
[1] 0

$message
[1] "Rcgmin seems to have converged"

$bdmsk
[1] 0 0 0

  > alvmma <- Rvmmmin(start1, hobbs.f, hobbs.g, bdmsk=c(0,0,0))
  > alvmma

$par
b1 b2 b3
 1  1  1

$value
[1] 23521

$counts
function gradient
      1          1

$convergence
[1] 0

$message
[1] "Rvmmminb appears to have converged"

$bdmsk
[1] 0 0 0

  > # Bounds and mask
  > ansmbm2 <- nlfb(startf1, hobbs.res, hobbs.jac, maskidx=c(2),
+                 lower=c(0,0,0), upper=c(200, 60, .3))
  > ansmbm2

nlmrt class object: x
residual sumsquares = 6181.2 on 12 observations
after 17 Jacobian and 28 function evaluations

```

name	coeff	SE	tstat	pval	gradient	JSingval
b1	50.4016	NA	NA	NA	0.0004618	162.2
b2	1 M	NA	NA	NA	0	0.4918
b3	0.198618	NA	NA	NA	-0.0746	0

```

  > a1cgbm2 <- Rcgmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,0,1),
+                 lower=c(0,0,0), upper=c(200, 60, .3))
  > a1cgbm2

$par
      b1      b2      b3
50.40179 1.00000 0.19861

$value
[1] 6181.2

$counts
[1] 1004 118

$convergence
[1] 1

$message
[1] "Too many function evaluations (> 1000) "

$bdmsk
[1] 1 0 1

```

```

> a1vmbm2 <- Rvmmmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,0,1),
+                   lower=c(0,0,0), upper=c(200, 60, .3))
> a1vmbm2

$par
      b1      b2      b3
50.40179  1.00000  0.19861

$value
[1] 6181.2

$counts
function gradient
      31      12

$convergence
[1] 0

$message
[1] "Rvmmminb appears to have converged"

$bdmsk
[1] 1 0 1

> # Active bound
> a1cgm2x <- Rcgmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,0,1),
+                 lower=c(0,0,0), upper=c(48, 60, .3))
> a1cgm2x

$par
      b1      b2      b3
48.00000  1.00000  0.21597

$value
[1] 6206.1

$counts
[1] 1005 115

$convergence
[1] 1

$message
[1] "Too many function evaluations (> 1000) "

$bdmsk
[1] -1 0 1

> a1vmm2x <- Rvmmmin(start1, hobbs.f, hobbs.g, bdmsk=c(1,0,1),
+                   lower=c(0,0,0), upper=c(48, 60, .3))
> a1vmm2x

$par
      b1      b2      b3
48.0  1.0  0.3

$value
[1] 6463.3

$counts
function gradient
      2      2

$convergence
[1] 2

$message
[1] "Rvmmminb appears to have converged"

```

```
$bdmsk
[1] -1  0 -1
```

## 6 Brief example of minpack.lm

Recently Kate Mullen provided some capability for the package `minpack.lm` to include bounds constraints. I am particularly happy that this effort is proceeding, as there are significant differences in how `minpack.lm` and `nlmrt` are built and implemented. They can be expected to have different performance characteristics on different problems. A lively dialogue between developers, and the opportunity to compare and check results can only improve the tools.

The examples below are a very quick attempt to show how to run the Ratkowsky-Huet problem with `nls.lm` from `minpack.lm`.

```
> require(minpack.lm)
> anlslm <- nls.lm(ones, lower=rep(-1000,4), upper=rep(1000,4), jres, jjac, yield=pastured)
> cat("anlslm from ones\n")
anlslm from ones
> print(strwrap(anlslm))
[1] "c(NaN, NaN, NaN, NaN)"
[2] "c(NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN)"
[3] "NaN, NaN, NaN, NaN, NaN, NaN, NaN"
[4] "c(NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN)"
[5] "4"
[6] "The cosine of the angle between 'fvec' and any column"
[7] "of the Jacobian is at most 'gtol' in absolute value."
[8] "list(t1 = 3, t2 = 2.3723939879224e-11, t3 ="
[9] "5.8039519205899e-10, t4 = 1.27525858056086e-09)"
[10] "3"
[11] "c(17533.3402000004, 16864.5616372991, NaN,"
[12] "3.95252516672997e-323)"
[13] "NaN"
> anlslmh <- nls.lm(huetstart, lower=rep(-1000,4), upper=rep(1000,4), jres, jjac, yield=pastured)
> cat("anlslmh from huetstart\n")
anlslmh from huetstart
> print(strwrap(anlslmh))
[1] "c(69.9551973916736, 61.6814877170941,"
[2] "-9.20891880263443, 2.37781455978467)"
[3] "c(9, -4.54037977686007, 105.318033221555,"
[4] "403.043210394646, -4.54037977686007,"
[5] "3.51002837648689, -39.5314537948583,"
[6] "-137.559566823766, 105.318033221555,"
[7] "-39.5314537948583, 1668.11894086464,"
[8] "6495.67702199831, 403.043210394646,"
[9] "-137.559566823766, 6495.67702199831,"
[10] "25481.4530263827)"
[11] "c(0.480682793156291, 0.669303022602289,"
[12] "-2.28431914156848, 0.843754801653787,"
[13] "0.734587578832198, 0.0665510313004489,"
[14] "-0.985814877917491, -0.0250630130722556,"
[15] "0.500317790294602)"
[16] "1"
[17] "Relative error in the sum of squares is at most"
[18] "'ftol'."
```



```

[19] "list(t1 = 3, t2 = 2.35105755434962, t3 ="
[20] "231.250186433367, t4 = 834.778914353851)"
[21] "42"
[22] "c(13386.9099465603, 13365.3097414383,"
[23] "13351.1970260154, 13321.6478455192, 13260.1135652244,"
[24] "13133.6391318145, 12877.8542053848, 12373.5432344283,"
[25] "11428.8257706578, 9832.87890178625, 7138.12187613237,"
[26] "3904.51162830831, 2286.64875980737, 1978.18149980306,"
[27] "1620.89081508973, 1140.58638304326, 775.173148616758,"
[28] "635.256627921479, 383.73614705125, 309.341249993346,"
[29] "219.735856060244, 177.398738179149, 156.718991828473,"
[30] "135.51359456819, 93.4016394568234, 72.8219383036211,"
[31] "66.3315609834918, 56.2809616213409, 54.9453021619838,"
[32] "53.6227655715768, 51.9760950696957, 50.1418078879665,"
[33] "48.1307021647518, 44.7097757109306, 42.8838792615115,"
[34] "32.3474231559263, 26.5253835687508, 15.352821554109,"
[35] "14.7215507012923, 8.37980617628203, 8.37589765770215,"
[36] "8.3758836534811, 8.37588355972578)"
[37] "8.37588355972578"

```

## References

- [1] Timur V. Elzhov, Katharine M. Mullen, Andrej-Nikolai Spiess, and Ben Bolker, *minpack.lm: R interface to the levenberg-marquardt nonlinear least-squares algorithm found in minpack, plus support for bounds*, R Project for Statistical Computing, 2012, R package version 1.1-6.
- [2] S. (Sylvie) Huet et al., *Statistical tools for nonlinear regression: a practical guide with S-PLUS examples*, Springer series in statistics, 1996.
- [3] J. J. Moré, B. S. Garbow, and K. E. Hillstom, *ANL-80-74, User Guide for MINPACK-1*, Tech. report, 1980.
- [4] J. C. Nash, *Compact numerical methods for computers : linear algebra and function minimisation*, Hilger, Bristol :, 1979 (English).
- [5] David A. Ratkowsky, *Nonlinear regression modeling: A unified practical approach*, Marcel Dekker Inc., New York and Basel, 1983.