

nlmrt-vignette

John C. Nash

August 5, 2012

Background

This vignette discusses the R package `nlmrt`, that aims to provide computationally robust tools for nonlinear least squares problems. Note that R already has the `nls()` function to solve nonlinear least squares problems, and this function has a large repertoire of tools for such problems. However, it is specifically NOT indicated for problems where the residuals are small or zero. Furthermore, it frequently fails to find a solution if starting parameters are provided that are not close enough to a solution. The tools of `nlmrt` are very much intended to cope with both these issues.

`nlmrt` tools generally do not return the large `nls`-style object. However, we do provide a tool `wrapnls` that will run either `nlxb` followed by a call to `nls`. The call to `nls` is adjusted to use the `port` algorithm if there are bounds constraints.

1 An example problem and its solution

Let us try an example initially presented by (Ratkowsky 1983) and developed by (Huet et al. 1996). This is a model for the regrowth of pasture. We set up the computation by putting the data for the problem in a data frame, and specifying the formula for the model. This can be as a formula object, but I have found that saving it as a character string seems to give fewer difficulties. Note the `" "` that implies "is modeled by". There must be such an element in the formula for this package (and for `nls()`). We also specify two sets of starting parameters, that is, the `ones` which is a trivial (but possibly unsuitable) start with all parameters set to 1, and `huetstart` which was suggested in (Huet et al. 1996). Finally we load the routines in the package `nlmrt`.

```
options(width = 60)
pastured <- data.frame(time = c(9, 14, 21, 28, 42, 57, 63, 70,
79),
  yield = c(8.93, 10.8, 18.59, 22.33, 39.35, 56.11, 61.73,
64.62, 67.08))
regmod <- "yield ~ t1 - t2*exp(-exp(t3+t4*log(time)))"
```

```

ones <- c(t1 = 1, t2 = 1, t3 = 1, t4 = 1) # all ones start
huetstart <- c(t1 = 70, t2 = 60, t3 = 0, t4 = 1)
require(nlmrt)

## Loading required package: nlmrt

```

Let us now call the routine `nlsmnqb` (even though we are not specifying bounds). We try both starts.

```

anmrt <- nlxb(regmod, start = ones, trace = FALSE,
              data = pastured)
print(anmrt)

## $resid
## [1] 0.48070 0.66931 -2.28433 0.84374 0.73458 0.06655
## [7] -0.98581 -0.02506 0.50032
##
## $jacobian
##      t1      t2      t3      t4
## [1,] 1 -0.9816 1.126 2.475
## [2,] 1 -0.9482 3.111 8.211
## [3,] 1 -0.8698 7.485 22.787
## [4,] 1 -0.7584 12.935 43.102
## [5,] 1 -0.4843 21.659 80.956
## [6,] 1 -0.2234 20.652 83.498
## [7,] 1 -0.1493 17.515 72.569
## [8,] 1 -0.0869 13.095 55.634
## [9,] 1 -0.0385 7.735 33.798
##
## $feval
## [1] 76
##
## $jeval
## [1] 50
##
## $coeffs
## [1] 69.955 61.681 -9.209 2.378
##
## $ssquares
## [1] 8.376
##

```

```

anmrtn <- try(nlxb(regmod, start = huetstart, trace = FALSE,
                  data = pastured))
print(strwrap(anmrtn))

```

```
## [1] "c(0.480699476110992, 0.669309701586503,"
## [2] "-2.28432650017661, 0.843738460841614,"
## [3] "0.734575256138093, 0.0665546618861583,"
## [4] "-0.985808933151056, -0.0250584603521418,"
## [5] "0.500316337120296)"
## [6] "c(1, 1, 1, 1, 1, 1, 1, 1, 1, -0.981567160420883,"
## [7] "-0.948192289406167, -0.869783557170751,"
## [8] "-0.758436212560273, -0.484272123696113,"
## [9] "-0.223383622127412, -0.149331587423979,"
## [10] "-0.0869019449646661, -0.0385020596618461,"
## [11] "1.12642043233262, 3.11132895498809, 7.48468988716119,"
## [12] "12.9349083313689, 21.6594224095687, 20.652293670436,"
## [13] "17.51548586967, 13.0949252904654, 7.73503096811733,"
## [14] "2.47499865833493, 8.2109754835055, 22.7873063008638,"
## [15] "43.1017598804902, 80.9557650898109, 83.4982821079476,"
## [16] "72.56901775625, 55.6337277915341, 33.7978144524062)"
## [17] "61"
## [18] "39"
## [19] "c(69.9551789601637, 61.6814436396711,"
## [20] "-9.20893535565824, 2.37781880027694)"
## [21] "8.37588355893792"
```

Note that the standard `nls()` of R fails to find a solution from either start.

```
anls <- try(nls(regmod, start = ones, trace = FALSE,
               data = pastured))
print(strwrap(anls))

## [1] "Error in nlsModel(formula, mf, start, wts) : singular"
## [2] "gradient matrix at initial parameter estimates"
```

```
anlsx <- try(nls(regmod, start = huetstart, trace = FALSE,
                 data = pastured))
print(strwrap(anlsx))

## [1] "Error in nls(regmod, start = huetstart, trace ="
## [2] "FALSE, data = pastured) : singular gradient"
```

In both cases, the `nls()` failed with a 'singular gradient'. This implies the Jacobian is effectively singular at some point. The Levenberg-Marquardt stabilization used in `nlxb` avoids this particular issue by augmenting the Jacobian until it is non-singular. The details of this common approach may be found elsewhere (Nash 1979). ?? Do we want a page ref?

There are some other tools for R that aim to solve nonlinear least squares problems. We have not yet been able to successfully use the INRA package

`nls2`. This is a quite complicated package and is not installable as a regular R package using `install.packages()`. Note that there is a very different package by the same name on CRAN by Gabor Grothendieck.

2 The `nls` solution

We can call `nls` after getting a potential nonlinear least squares solution using `nlxb`. Package `nlmrt` has function `wrapnls` to allow this to be carried out automatically. Thus,

```
awnls <- wrapnls(regmod, start = ones, data = pastured)
print(awnls)

## Nonlinear regression model
##   model:  yield ~ t1 - t2 * exp(-exp(t3 + t4 * log(time)))
##   data:  data
##    t1    t2    t3    t4
## 69.96 61.68 -9.21  2.38
## residual sum-of-squares: 8.38
##
## Number of iterations to convergence: 0
## Achieved convergence tolerance: 8.33e-08

cat("Note that the above is just the nls() summary result.\n")

## Note that the above is just the nls() summary result.
```

3 Problems specified by residual functions

The model expressions in R, such as

```
yield ~ t1 - t2*exp(-exp(t3+t4*log(time)))
```

are an extremely helpful feature of the language. Moreover, they are used to compute symbolic or automatic derivatives, so we do not have to rely on numerical approximations for the Jacobian of the nonlinear least squares problem. However, there are many situations where the expression structure is not flexible enough to allow us to define our residuals, or where the construction of the residuals is simply too complicated. In such cases it is helpful to have tools that work with R functions.

Once we have an R function for the residuals, we can use the safeguarded Marquardt routine `nlfb` from package `nlmrt` or else the routine `nls.lm` from package `minpack.lm` (Elzhov, Mullen, Spiess, and Bolker 2012). The latter is built on the Minpack Fortran codes of (Moré, Garbow, and Hillstom 1980) implemented by Kate Mullen. `nlfb` is written entirely in R, and is intended

to be quite aggressive in ensuring it finds a good minimum. Thus these two approaches have somewhat different characteristics.

Let us consider a slightly different problem, called WEEDS. Here the objective is to model a set of 12 data points (density y of weeds at annual time points tt) versus the time index. (A minor note: use of t rather than tt in R may encourage confusion with the transpose function $t()$, so I tend to avoid plain t .) The model suggested was a 3-parameter logistic function,

$$y_{model} = b_1 / (1 + b_2 \exp(-b_3 tt))$$

and while it is possible to use this formulation, a scaled version gives slightly better results

$$y_{model} = 100b_1 / (1 + 10b_2 \exp(-0.1b_3 tt))$$

The residuals for this latter model (in form "model" minus "data") are coded in R in the following code chunk in the function `shobbs.res`. We have also coded the Jacobian for this model as `shobbs.jac`

```
shobbs.res <- function(x) {
  # scaled Hobbs weeds problem - residual
  # This variant uses looping
  if (length(x) != 3)
    stop("hobbs.res - parameter vector n!=3")
  y <- c(5.308, 7.24, 9.638, 12.866, 17.069, 23.192, 31.443,
        38.558, 50.156, 62.948, 75.995, 91.972)
  tt <- 1:12
  res <- 100 * x[1] / (1 + x[2] * 10 * exp(-0.1 * x[3] * tt)) -
    y
}

shobbs.jac <- function(x) {
  # scaled Hobbs weeds problem - Jacobian
  jj <- matrix(0, 12, 3)
  tt <- 1:12
  yy <- exp(-0.1 * x[3] * tt) # We don't need data for the
  # Jacobian
  zz <- 100 / (1 + 10 * x[2] * yy)
  jj[tt, 1] <- zz
  jj[tt, 2] <- -0.1 * x[1] * zz * zz * yy
  jj[tt, 3] <- 0.01 * x[1] * zz * zz * yy * x[2] * tt
  return(jj)
}
```

With package `nlmrt`, function `nlfb` can be used to estimate the parameters of the WEEDS problem as follows, where we use the naive starting point where all parameters are 1.

```

st <- c(b1 = 1, b2 = 1, b3 = 1)
ans1 <- nlfb(st, shobbs.res, shobbs.jac, trace = FALSE)
print(ans1)

## $resid
## [1] 0.01190 -0.03276 0.09203 0.20878 0.39263 -0.05759
## [7] -1.10573 0.71579 -0.10765 -0.34840 0.65259 -0.28757
##
## $jacobian
##      [,1] [,2] [,3]
## [1,] 2.712 -1.054 0.5176
## [2,] 3.674 -1.414 1.3885
## [3,] 4.960 -1.884 2.7742
## [4,] 6.664 -2.486 4.8814
## [5,] 8.901 -3.240 7.9537
## [6,] 11.792 -4.157 12.2438
## [7,] 15.464 -5.224 17.9522
## [8,] 20.019 -6.399 25.1294
## [9,] 25.511 -7.594 33.5526
## [10,] 31.908 -8.683 42.6252
## [11,] 39.069 -9.513 51.3725
## [12,] 46.733 -9.948 58.6047
##
## $feval
## [1] 24
##
## $jeval
## [1] 15
##
## $coeffs
## [1] 1.962 4.909 3.136
##
## $ssquares
## [1] 2.587
##

```

This works very well, with almost identical iterates as given by `nlxb`. (Since the algorithms are the same, this should be the case.) Note that we turn off the `trace` output. There is also the possibility of interrupting the iterations to `watch` the progress. Changing the value of `watch` in the call to `nlfb` below allows this. In this code chunk, we use an internal numerical approximation to the Jacobian.

```

cat("No jacobian function - use internal approximation\n")

## No jacobian function -- use internal approximation

```

```

ans1n <- nlfb(st, shobbs.res, trace = FALSE, control = list(watch
= FALSE)) # NO jacfn
print(ans1n)

## $resid
## [1] 0.01190 -0.03276 0.09203 0.20878 0.39263 -0.05759
## [7] -1.10573 0.71579 -0.10765 -0.34840 0.65259 -0.28757
##
## $jacobian
##      [,1] [,2] [,3]
## [1,] 2.712 -1.054 0.5176
## [2,] 3.674 -1.414 1.3885
## [3,] 4.960 -1.884 2.7742
## [4,] 6.664 -2.486 4.8814
## [5,] 8.901 -3.240 7.9537
## [6,] 11.792 -4.157 12.2438
## [7,] 15.464 -5.224 17.9522
## [8,] 20.019 -6.399 25.1294
## [9,] 25.511 -7.594 33.5526
## [10,] 31.908 -8.683 42.6252
## [11,] 39.069 -9.513 51.3725
## [12,] 46.733 -9.948 58.6047
##
## $feval
## [1] 29
##
## $jeval
## [1] 15
##
## $coeffs
## [1] 1.962 4.909 3.136
##
## $ssquares
## [1] 2.587
##

```

Note that we could also form the sum of squares function and the gradient and use a function minimization code. The next code block shows how this is done, creating the sum of squares function and its gradient, then using the `optimx` package to call a number of minimizers simultaneously.

```

shobbs.f <- function(x) {
  res <- shobbs.res(x)
  as.numeric(crossprod(res))
}
shobbs.g <- function(x) {

```

```

    res <- shobbs.res(x) # This is NOT efficient - we generally
have res already calculated
    JJ <- shobbs.jac(x)
    2 * as.vector(crossprod(JJ, res))
}
require(optimx)

## Loading required package: optimx
## Loading required package: numDeriv

aopx <- optimx(st, shobbs.f, shobbs.g, control = list(all.methods
= TRUE))

## Attaching package: 'RvmminCRAN'
## The following object(s) are masked from 'package:optimx':
##
## optansout
## Loading required package: methods
## end topstuff in optimxCRAN

optansout(aopx, NULL) # no file output

##
## par
## 2 1.912, 4.825, 3.159
## 3 1.964, 4.912, 3.134
## 7 1.962, 4.909, 3.136
## 5 1.962, 4.909, 3.136
## 1 1.962, 4.909, 3.136
## 12 1.962, 4.909, 3.136
## 11 1.962, 4.909, 3.136
## 4 1.962, 4.909, 3.136
## 10 1.962, 4.909, 3.136
## 6 1.962, 4.909, 3.136
## 9 1.962, 4.909, 3.136
## 8 1.962, 4.909, 3.136
## fvalues method fns grs itns conv KKT1 KKT2 xtimes
## 2 2.668 CG 427 101 NULL 1 FALSE TRUE 0.012
## 3 2.588 Nelder-Mead 196 NA NULL 0 FALSE TRUE 0.004
## 7 2.587 spg 188 NA 150 0 TRUE TRUE 0.032
## 5 2.587 nlm NA NA 50 0 TRUE TRUE 0.004
## 1 2.587 BFGS 119 36 NULL 0 TRUE TRUE 0.004
## 12 2.587 bobyqa 705 NA NULL 0 TRUE TRUE 0.02
## 11 2.587 newuoa 1957 NA NULL 0 TRUE TRUE 0.056
## 4 2.587 L-BFGS-B 41 41 NULL 0 TRUE TRUE 0.004
## 10 2.587 Rvmmin 83 47 NULL 0 TRUE TRUE 0.012
## 6 2.587 nlminb 31 29 28 0 TRUE TRUE 0.004

```



```

## 9      2.587      Rcgmin 138 50 NULL 0 TRUE TRUE 0.008
## 8      2.587      ucminf 46 46 NULL 0 TRUE TRUE 0.004
## [1] TRUE

cat("\nNow with numerical gradient approximation or derivative
free methods\n")

##
## Now with numerical gradient approximation or derivative free methods

aopxn <- optimx(st, shobbs.f, control = list(all.methods = TRUE))

## end topstuff in optimxCRA
## Warning: A NULL gradient function is being replaced with
## fwd diff for Rcgmin
## function(x) {
##     res <- shobbs.res(x)
##     as.numeric(crossprod(res))
## }
## Warning: Numerical gradients may be inappropriate for
## Rvmmmin

optansout(aopxn, NULL) # no file output

##
## par
## 2 1.800, 4.597, 3.208
## 3 1.964, 4.912, 3.134
## 8 1.962, 4.909, 3.136
## 7 1.962, 4.909, 3.136
## 1 1.962, 4.909, 3.136
## 10 1.962, 4.909, 3.136
## 4 1.962, 4.909, 3.136
## 5 1.962, 4.909, 3.136
## 12 1.962, 4.909, 3.136
## 11 1.962, 4.909, 3.136
## 9 1.962, 4.909, 3.136
## 6 1.962, 4.909, 3.136
## fvalues      method fns grs itns conv KKT1 KKT2 xtimes
## 2      3.83      CG 413 101 NULL 1 FALSE TRUE 0.02
## 3      2.588 Nelder-Mead 196 NA NULL 0 FALSE TRUE 0.004
## 8      2.587      ucminf 45 45 NULL 0 FALSE TRUE 0.004
## 7      2.587      spg 174 NA 135 0 TRUE TRUE 0.032
## 1      2.587      BFGS 118 36 NULL 0 TRUE TRUE 0.004
## 10     2.587      Rvmmmin 83 44 NULL 0 TRUE TRUE 0.016
## 4      2.587      L-BFGS-B 45 45 NULL 0 TRUE TRUE 0.004
## 5      2.587      nlm NA NA 50 0 TRUE TRUE 0.004
## 12     2.587      bobyqa 705 NA NULL 0 TRUE TRUE 0.02

```

```
## 11 2.587 newuoa 1957 NA NULL 0 TRUE TRUE 0.056
## 9 2.587 Rcgmin 128 48 NULL 0 TRUE TRUE 0.064
## 6 2.587 nlminb 32 93 27 0 TRUE TRUE 0.004
## [1] TRUE
```

We see that most of the minimizers work with either the analytic or approximated gradient. The 'CG' option of function `optim()` does not do very well in either case. As the author of the original step and description and then Turbo Pascal code, I can say I was never very happy with this method and replaced it recently with `Rcgmin` from the package of the same name, in the process adding the possibility of bounds or masks constraints.

4 Converting an expression to a function

Clearly if we have an expression, it would be nice to be able to automatically convert this to a function, if possible also getting the derivatives. Indeed, it is possible to convert an expression to a function, and there are several ways to do this (references??). In package `nlmrt` we provide the tools `model2grfun.R`, `model2jacfun.R`, `model2resfun.R`, and `model2ssfun.R` to convert a model expression to a function to compute the gradient, Jacobian, residuals or sum of squares functions respectively. We do not provide any tool for converting a function for the residuals back to an expression, as functions can use structures that are not easily expressed as R expressions.

Below are code chunks to illustrate the generation of the residual, sum of squares, Jacobian and gradient code for the Ratkowsky problem used earlier in the vignette. The commented-out first line shows how we would use one of these function generators to output the function to a file named "testresfn.R". However, it is not necessary to generate the file.

First, let us generate the residuals. We must supply the names of the parameters, and do this via the starting vector of parameters `ones`. The actual values are not needed by `model2resfun`, just the names. Other names are drawn from the variables used in the model expression `regmod`.

```
# jres<-model2resfun(regmod, ones, funname='myxres',
# file='testresfn.R')
jres <- model2resfun(regmod, ones)
print(jres)

## function (prm, yield = NULL, time = NULL)
## {
##   t1 <- prm[[1]]
##   t2 <- prm[[2]]
##   t3 <- prm[[3]]
##   t4 <- prm[[4]]
##   resids <- as.numeric(eval(t1 - t2 * exp(-exp(t3 + t4 * log(time)))) -
```

```

##         yield))
## }
## <environment: 0x9ee55b4>

valjres <- jres(ones, yield = pastured$yield, time =
pastured$time)
cat("valjres:")

## valjres:

print(valjres)

## [1] -7.93 -9.80 -17.59 -21.33 -38.35 -55.11 -60.73 -63.62
## [9] -66.08

```

Now let us also generate the Jacobian and test it using the numerical approximations from package `numDeriv`.

```

jjac <- model2jacfun(regmod, ones)
print(jjac)

## function (prm, yield = NULL, time = NULL)
## {
##     t1 <- prm[[1]]
##     t2 <- prm[[2]]
##     t3 <- prm[[3]]
##     t4 <- prm[[4]]
##     localdf <- data.frame(yield, time)
##     jstruc <- with(localdf, eval({
##         .expr1 <- log(time)
##         .expr4 <- exp(t3 + t4 * .expr1)
##         .expr6 <- exp(-.expr4)
##         .value <- t1 - t2 * .expr6 - yield
##         .grad <- array(0, c(length(.value), 4), list(NULL, c("t1",
##             "t2", "t3", "t4")))
##         .grad[, "t1"] <- 1
##         .grad[, "t2"] <- -.expr6
##         .grad[, "t3"] <- t2 * (.expr6 * .expr4)
##         .grad[, "t4"] <- t2 * (.expr6 * (.expr4 * .expr1))
##         attr(.value, "gradient") <- .grad
##         .value
##     })))
##     jacmat <- attr(jstruc, "gradient")
##     return(jacmat)
## }
## <environment: 0x99dbbb0>

```

```

# Note that we now need some data!
valjjac <- jjac(ones, yield = pastured$yield, time =
pastured$time)
cat("valjac:")

## valjac:

print(valjjac)

##          t1          t2          t3          t4
## [1,]  1 -2.372e-11  5.804e-10  1.275e-09
## [2,]  1 -2.968e-17  1.130e-15  2.981e-15
## [3,]  1 -1.617e-25  9.232e-24  2.811e-23
## [4,]  1 -8.811e-34  6.706e-32  2.235e-31
## [5,]  1 -2.615e-50  2.986e-48  1.116e-47
## [6,]  1 -5.123e-68  7.938e-66  3.209e-65
## [7,]  1 -4.230e-75  7.243e-73  3.001e-72
## [8,]  1 -2.304e-83  4.385e-81  1.863e-80
## [9,]  1 -5.467e-94  1.174e-91  5.130e-91

# Now compute the numerical approximation
Jn <- jacobian(jres, ones, , yield = pastured$yield,
time = pastured$time)
cat("maxabsdiff=", max(abs(Jn - valjjac)), "\n")

## maxabsdiff= 3.774e-10

```

As with the WEEDS problem, we can compute the sum of squares function and the gradient.

```

ssfn <- model2ssfun(regmod, ones) # problem getting the data
attached!
print(ssfn)

## function (prm, yield = NULL, time = NULL)
## {
##   t1 <- prm[[1]]
##   t2 <- prm[[2]]
##   t3 <- prm[[3]]
##   t4 <- prm[[4]]
##   resids <- as.numeric(eval(t1 - t2 * exp(-exp(t3 + t4 * log(time))) -
##   yield))
##   ss <- as.numeric(crossprod(resids))
## }
## <environment: 0xad8728>

valss <- ssfn(ones, yield = pastured$yield, time = pastured$time)

```

```

cat("valss: ", valss, "\n")
## valss: 17533

grfn <- model2grfun(regmod, ones) # problem getting the data
attached!
print(grfn)

## function (prm, yield = NULL, time = NULL)
## {
##   t1 <- prm[[1]]
##   t2 <- prm[[2]]
##   t3 <- prm[[3]]
##   t4 <- prm[[4]]
##   localdf <- data.frame(yield, time)
##   jstruc <- with(localdf, eval({
##     .expr1 <- log(time)
##     .expr4 <- exp(t3 + t4 * .expr1)
##     .expr6 <- exp(-.expr4)
##     .value <- t1 - t2 * .expr6 - yield
##     .grad <- array(0, c(length(.value), 4), list(NULL, c("t1",
##       "t2", "t3", "t4")))
##     .grad[, "t1"] <- 1
##     .grad[, "t2"] <- -.expr6
##     .grad[, "t3"] <- t2 * (.expr6 * .expr4)
##     .grad[, "t4"] <- t2 * (.expr6 * (.expr4 * .expr1))
##     attr(.value, "gradient") <- .grad
##     .value
##   })))
##   jacmat <- attr(jstruc, "gradient")
##   resids <- as.numeric(eval(t1 - t2 * exp(-exp(t3 + t4 * log(time)))) -
##     yield))
##   grj <- as.vector(2 * crossprod(jacmat, resids))
## }
## <environment: 0x9f2b6b4>

valgr <- grfn(ones, yield = pastured$yield, time = pastured$time)
cat("valgr:")
## valgr:
print(valgr)
## [1] -6.811e+02 3.763e-10 -9.205e-09 -2.023e-08

gn <- grad(ssfn, ones, yield = pastured$yield, time =
pastured$time)
cat("maxabsdiff=", max(abs(gn - valgr)), "\n")
## maxabsdiff= 7.477e-08

```

Moreover, we can use the Huet starting parameters as a double check on our conversion of the expression to various optimization-style functions.

```
cat("\n\nHuetstart:")

##
##
## Huetstart:

print(huetstart)

## t1 t2 t3 t4
## 70 60 0 1

valjres <- jres(huetstart, yield = pastured$yield,
               time = pastured$time)
cat("valjres:")

## valjres:

print(valjres)

## [1] 61.06 59.20 51.41 47.67 30.65 13.89 8.27 5.38 2.92

valss <- ssfn(huetstart, yield = pastured$yield, time =
pastured$time)
cat("valss:", valss, "\n")

## valss: 13387

valjjac <- jjac(huetstart, yield = pastured$yield,
               time = pastured$time)
cat("valjac:")

## valjac:

print(valjjac)

##          t1          t2          t3          t4
## [1,]  1 -1.234e-04  6.664e-02  1.464e-01
## [2,]  1 -8.315e-07  6.985e-04  1.843e-03
## [3,]  1 -7.583e-10  9.554e-07  2.909e-06
## [4,]  1 -6.914e-13  1.162e-09  3.871e-09
## [5,]  1 -5.750e-19  1.449e-15  5.415e-15
## [6,]  1 -1.759e-25  6.015e-22  2.432e-21
## [7,]  1 -4.360e-28  1.648e-24  6.828e-24
## [8,]  1 -3.975e-31  1.670e-27  7.094e-27
## [9,]  1 -4.906e-35  2.325e-31  1.016e-30
```

```

Jn <- jacobian(jres, huetstart, , yield = pastured$yield,
               time = pastured$time)
cat("maxabsdiff=", max(abs(Jn - valjjac)), "\n")

## maxabsdiff= 5.395e-10

valgr <- grfn(huetstart, yield = pastured$yield, time =
pastured$time)
cat("valgr:")

## valgr:

print(valgr)

## [1] 560.90509 -0.01517 8.22138 18.10084

gn <- grad(ssfn, huetstart, yield = pastured$yield,
           time = pastured$time)
cat("maxabsdiff=", max(abs(gn - valgr)), "\n")

## maxabsdiff= 5.953e-08

```

Now that we have these functions, let us apply them with `nlfb`.

```

cat("All ones to start\n")

## All ones to start

anlfb <- nlfb(ones, jres, jjac, trace = FALSE, yield =
pastured$yield,
             time = pastured$time)
print(strwrap(anlfb))

## [1] "c(0.480699475409779, 0.669309701325741,"
## [2] "-2.28432649983562, 0.843738461541676,"
## [3] "0.734575256578069, 0.0665546616416748,"
## [4] "-0.985808933450038, -0.0250584605193325,"
## [5] "0.500316337308163)"
## [6] "c(1, 1, 1, 1, 1, 1, 1, 1, 1, -0.981567160415026,"
## [7] "-0.948192289394349, -0.869783557151951,"
## [8] "-0.758436212539591, -0.484272123689345,"
## [9] "-0.22338362214097, -0.14933158744104,"
## [10] "-0.086901944981799, -0.0385020596749348,"
## [11] "1.12642043272705, 3.1113289557883, 7.48468988842378,"
## [12] "12.9349083327494, 21.6594224104496, 20.6522936715837,"
## [13] "17.5154858712384, 13.0949252924535, 7.73503097021314,"
## [14] "2.47499865920158, 8.21097548561731, 22.7873063047078,"

```

```
## [15] "43.1017598850905, 80.9557650931036, 83.498282112588,"
## [16] "72.569017762748, 55.6337277999807, 33.7978144615637)"
## [17] "74"
## [18] "48"
## [19] "c(69.9551789612429, 61.6814436418531,"
## [20] "-9.20893535490747, 2.37781880008123)"
## [21] "8.37588355893788"

cat("Huet start\n")

## Huet start

anlfbh <- nlfb(huetstart, jres, jjac, trace = FALSE,
  yield = pastured$yield, time = pastured$time)
print(strwrap(anlfbh))

## [1] "c(0.480699465869456, 0.669309697775223,"
## [2] "-2.28432649519877, 0.84373847107085,"
## [3] "0.734575262591456, 0.0665546583437617,"
## [4] "-0.985808937499776, -0.0250584627932966,"
## [5] "0.500316339841277)"
## [6] "c(1, 1, 1, 1, 1, 1, 1, 1, 1, -0.981567160335378,"
## [7] "-0.94819228923362, -0.869783556896137,"
## [8] "-0.75843621225793, -0.484272123596337,"
## [9] "-0.223383622324199, -0.149331587672017,"
## [10] "-0.0869019452139657, -0.0385020598524092,"
## [11] "1.12642043808933, 3.11132896666899, 7.48468990559557,"
## [12] "12.9349083515304, 21.6594224224275, 20.652293687139,"
## [13] "17.5154858924942, 13.0949253194057, 7.73503099863509,"
## [14] "2.47499867098372, 8.21097551433206, 22.7873063569877,"
## [15] "43.1017599476725, 80.9557651378729, 83.498282175479,"
## [16] "72.5690178508139, 55.6337279144867, 33.7978145857519)"
## [17] "60"
## [18] "37"
## [19] "c(69.9551789758633, 61.6814436714725,"
## [20] "-9.20893534470294, 2.37781879742191)"
## [21] "8.37588355893793"
```

5 Using bounds and masks

6 Brief comparison with minpack.lm


```

require(minpack.lm)

## Loading required package: minpack.lm

anls1m <- nls.lm(ones, lower = rep(-1000, 4), upper = rep(1000,
  4), jres, jjac, yield = pastured$yield, time = pastured$time)
cat("anls1m from ones\n")

## anls1m from ones

print(strwrap(anls1m))

## [1] "c(NaN, NaN, NaN, NaN)"
## [2] "c(NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN)"
## [3] "NaN, NaN, NaN, NaN, NaN, NaN, NaN)"
## [4] "c(NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN, NaN)"
## [5] "4"
## [6] "The cosine of the angle between `fvec' and any column"
## [7] "of the Jacobian is at most `gtol' in absolute value."
## [8] "list(t1 = 3, t2 = 2.3723939879224e-11, t3 ="
## [9] "5.8039519205899e-10, t4 = 1.27525858056086e-09)"
## [10] "3"
## [11] "c(17533.3402000004, 16864.5616372991, NaN,"
## [12] "1.112549661455e-308)"
## [13] "NaN"

anls1mh <- nls.lm(huetstart, lower = rep(-1000, 4),
  upper = rep(1000, 4), jres, jjac, yield = pastured$yield,
  time = pastured$time)
cat("anls1mh from huetstart\n")

## anls1mh from huetstart

print(strwrap(anls1mh))

## [1] "c(69.9551973916736, 61.6814877170941,"
## [2] "-9.20891880263443, 2.37781455978467)"
## [3] "c(9, -4.54037977686007, 105.318033221555,"
## [4] "403.043210394647, -4.54037977686007,"
## [5] "3.51002837648689, -39.5314537948583,"
## [6] "-137.559566823766, 105.318033221555,"
## [7] "-39.5314537948583, 1668.11894086464,"
## [8] "6495.67702199832, 403.043210394647,"
## [9] "-137.559566823766, 6495.67702199832,"
## [10] "25481.4530263827)"
## [11] "c(0.480682793156298, 0.669303022602289,"
## [12] "-2.28431914156848, 0.84375480165378,"

```

```
## [13] "0.734587578832198, 0.0665510313004845,"
## [14] "-0.985814877917491, -0.0250630130722556,"
## [15] "0.500317790294616)"
## [16] "1"
## [17] "Relative error in the sum of squares is at most"
## [18] "`ftol'."
## [19] "list(t1 = 3, t2 = 2.35105755434962, t3 ="
## [20] "231.250186433367, t4 = 834.778914353853)"
## [21] "42"
## [22] "c(13386.9099465603, 13365.3097414383,"
## [23] "13351.1970260154, 13321.6478455192, 13260.1135652244,"
## [24] "13133.6391318145, 12877.8542053848, 12373.5432344283,"
## [25] "11428.8257706578, 9832.87890178625, 7138.12187613238,"
## [26] "3904.51162830831, 2286.64875980737, 1978.18149980306,"
## [27] "1620.89081508973, 1140.58638304326, 775.173148616759,"
## [28] "635.256627921485, 383.73614705125, 309.34124999335,"
## [29] "219.735856060243, 177.39873817915, 156.718991828473,"
## [30] "135.513594568191, 93.4016394568244, 72.8219383036213,"
## [31] "66.331560983492, 56.2809616213412, 54.9453021619837,"
## [32] "53.6227655715772, 51.9760950696957, 50.1418078879664,"
## [33] "48.130702164752, 44.7097757109316, 42.8838792615125,"
## [34] "32.3474231559281, 26.5253835687528, 15.3528215541113,"
## [35] "14.7215507012991, 8.37980617628204, 8.37589765770224,"
## [36] "8.37588365348112, 8.37588355972579)"
## [37] "8.37588355972579"
```

References

- Elzhov, T. V., K. M. Mullen, A.-N. Spiess, and B. Bolker (2012). *minpack.lm: R interface to the Levenberg-Marquardt nonlinear least-squares algorithm found in MINPACK, plus support for bounds*. R Project for Statistical Computing. R package version 1.1-6.
- Huet, S. S. et al. (1996). *Statistical tools for nonlinear regression: a practical guide with S-PLUS examples*. Springer series in statistics.
- Moré, J. J., B. S. Garbow, and K. E. Hillstom (1980). ANL-80-74, User Guide for MINPACK-1. Technical report.
- Nash, J. C. (1979). *Compact Numerical Methods for Computers: Linear Algebra and Function Minimisation*. Bristol: Adam Hilger. Second Edition, 1990, Bristol: Institute of Physics Publications.
- Ratkowsky, D. A. (1983). *Nonlinear Regression Modeling: A Unified Practical Approach*. New York and Basel: Marcel Dekker Inc.